

✓ Task 1

```
from google.colab import drive
drive.mount('/content/drive')
```

↗ Mounted at /content/drive

✓ Improve the Previous Model

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)
```

```
test_datagen = ImageDataGenerator(rescale=1./255)
```

```
train_generator = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/AI-ML/week-5/FruitinAmazon/train',
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical'
)
```

```
validation_generator = test_datagen.flow_from_directory(
    '/content/drive/MyDrive/AI-ML/week-5/FruitinAmazon/test',
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical'
)
```

↗ Found 90 images belonging to 6 classes.
Found 30 images belonging to 6 classes.

✓ Build a Deeper CNN Model with Batch Normalization & Dropout

```

➡ /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` / `input_dim` argument to the layer constructor. Use the `build` method of the layer instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 32)	896
batch_normalization (BatchNormalization)	(None, 62, 62, 32)	128
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 29, 29, 64)	18,496
batch_normalization_1 (BatchNormalization)	(None, 29, 29, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 64)	0
dropout_3 (Dropout)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 12, 12, 128)	73,856
batch_normalization_2 (BatchNormalization)	(None, 12, 12, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
dropout_4 (Dropout)	(None, 6, 6, 128)	0
flatten (Flatten)	(None, 4608)	0
dense_3 (Dense)	(None, 256)	1,179,904
dropout_5 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 10)	2,570

Total params: 1,276,618 (4.87 MB)
Trainable params: 1,276,170 (4.87 MB)
Non-trainable params: 448 (1.75 KB)

- ✓ Transfer Learning with VGG16

```

from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D

# Load VGG16 without the top classification layer
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze base model layers
for layer in base_model.layers:
    layer.trainable = False

```

▼ Add Custom Layers

```

num_classes = len(train_generator.class_indices)

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1024, activation='relu')(x)
x = Dense(num_classes, activation='softmax')(x)
model = Model(inputs=base_model.input, outputs=x)

```

▼ Compile & Train the Model

```

from tensorflow.keras.optimizers import Adam

model.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])

model.fit(train_generator, epochs=10, validation_data=validation_generator)

```

→ /usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class self._warn_if_super_not_called()


```

Epoch 1/10
3/3 ————— 76s 25s/step - accuracy: 0.1915 - loss: 1.8709 - val_accuracy: 0.3000 - val_loss: 1.6741
Epoch 2/10
3/3 ————— 3s 1s/step - accuracy: 0.3508 - loss: 1.6312 - val_accuracy: 0.3667 - val_loss: 1.5466
Epoch 3/10
3/3 ————— 4s 595ms/step - accuracy: 0.4735 - loss: 1.4632 - val_accuracy: 0.5333 - val_loss: 1.4609
Epoch 4/10
3/3 ————— 2s 611ms/step - accuracy: 0.7506 - loss: 1.2629 - val_accuracy: 0.5000 - val_loss: 1.4334
Epoch 5/10
3/3 ————— 2s 687ms/step - accuracy: 0.8033 - loss: 1.1434 - val_accuracy: 0.4333 - val_loss: 1.3576
Epoch 6/10
3/3 ————— 2s 669ms/step - accuracy: 0.8002 - loss: 1.0428 - val_accuracy: 0.5667 - val_loss: 1.3070
Epoch 7/10
3/3 ————— 3s 978ms/step - accuracy: 0.7946 - loss: 0.9589 - val_accuracy: 0.5333 - val_loss: 1.3066
Epoch 8/10
3/3 ————— 2s 749ms/step - accuracy: 0.7992 - loss: 0.8814 - val_accuracy: 0.5000 - val_loss: 1.2932
Epoch 9/10
3/3 ————— 2s 614ms/step - accuracy: 0.7868 - loss: 0.7906 - val_accuracy: 0.5667 - val_loss: 1.2585
Epoch 10/10
3/3 ————— 2s 673ms/step - accuracy: 0.8738 - loss: 0.6777 - val_accuracy: 0.5667 - val_loss: 1.2525
<keras.src.callbacks.history.History at 0x7ecdd02659d0>

```

▼ Model Analysis

```
model.summary()
```

 Model: "functional_1"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
global_average_pooling2d_3 (GlobalAveragePooling2D)	(None, 512)	0
dense_5 (Dense)	(None, 1024)	525,312
dense_6 (Dense)	(None, 6)	6,150

Total params: 16,309,076 (62.21 MB)
 Trainable params: 531,462 (2.03 MB)
 Non-trainable params: 14,714,688 (56.13 MB)
 Optimizer params: 1,062,926 (4.05 MB)

Visualize Training Performance

```


import matplotlib.pyplot as plt

history = model.fit(train_generator, epochs=10, validation_data=validation_generator)

plt.plot(history.history['accuracy'], label='train_accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.legend()
plt.show()

```

Epoch 1/10
3/3 ————— 3s 689ms/step - accuracy: 0.8135 - loss: 0.6785 - val_accuracy: 0.6000 - val_loss: 1.2754
Epoch 2/10
3/3 ————— 2s 614ms/step - accuracy: 0.9270 - loss: 0.6272 - val_accuracy: 0.6000 - val_loss: 1.2428
Epoch 3/10
3/3 ————— 2s 673ms/step - accuracy: 0.8715 - loss: 0.6131 - val_accuracy: 0.6000 - val_loss: 1.1491
Epoch 4/10
3/3 ————— 2s 578ms/step - accuracy: 0.8732 - loss: 0.5267 - val_accuracy: 0.6000 - val_loss: 1.1786
Epoch 5/10
3/3 ————— 2s 622ms/step - accuracy: 0.8803 - loss: 0.5441 - val_accuracy: 0.6667 - val_loss: 1.2328
Epoch 6/10
3/3 ————— 3s 831ms/step - accuracy: 0.9015 - loss: 0.4733 - val_accuracy: 0.6667 - val_loss: 1.1418
Epoch 7/10
3/3 ————— 2s 671ms/step - accuracy: 0.9499 - loss: 0.4446 - val_accuracy: 0.6000 - val_loss: 1.1510
Epoch 8/10
3/3 ————— 2s 629ms/step - accuracy: 0.9315 - loss: 0.3983 - val_accuracy: 0.6667 - val_loss: 1.1114
Epoch 9/10
3/3 ————— 2s 654ms/step - accuracy: 0.9646 - loss: 0.3852 - val_accuracy: 0.7000 - val_loss: 1.0882
Epoch 10/10
3/3 ————— 2s 651ms/step - accuracy: 0.9750 - loss: 0.3385 - val_accuracy: 0.7000 - val_loss: 1.0619



Task 2

Data Preparation & Augmentation

```

from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_dir = '/content/drive/MyDrive/AI-ML/week-5/FruitinAmazon/train'
val_dir = '/content/drive/MyDrive/AI-ML/week-5/FruitinAmazon/test'

# Data augmentation
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    zoom_range=0.2,
    width_shift_range=0.2,
    height_shift_range=0.2,
    horizontal_flip=True
)


val_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical'
)

validation_generator = val_datagen.flow_from_directory(
    val_dir,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    shuffle=False # Important for inference output
)

num_classes = len(train_generator.class_indices)

```

 Found 90 images belonging to 6 classes.
Found 30 images belonging to 6 classes.

✓ Load VGG16 (Pre-trained) and Freeze Layers

```

from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout
from tensorflow.keras.optimizers import Adam

base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze all layers
for layer in base_model.layers:
    layer.trainable = False

```

✓ Add Custom Layers

```

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(512, activation='relu')(x)
x = Dropout(0.5)(x)
x = Dense(num_classes, activation='softmax')(x)

model = Model(inputs=base_model.input, outputs=x)

```

✓ Compile & Train the Model

```
model.compile(optimizer=Adam(learning_rate=0.0001), loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model.fit(  
    train_generator,  
    epochs=10,  
    validation_data=validation_generator  
)
```

↻ Epoch 1/10
3/3 ————— 8s 2s/step - accuracy: 0.1553 - loss: 2.0625 - val_accuracy: 0.1000 - val_loss: 1.8220
Epoch 2/10
3/3 ————— 3s 610ms/step - accuracy: 0.3104 - loss: 1.8707 - val_accuracy: 0.3000 - val_loss: 1.7073